

TMC, MP-RWS, MPIW, the Source Term Trick, and QEM

How alan does inference, plus a look at some newer empirical follow-ups

The problem

Tensor Monte Carlo

Massively-Parallel RWS

MPIW

The source term trick

QEM

How it fits together

Further work

The problem

Why one sample per latent isn't enough

- Standard importance sampling / VI: draw *one* sample of the whole latent vector z from Q , reweight by $p(z, x)/q(z)$.
- Better: draw K *whole-vector* samples (IWAE-style) and average: tighter bound, lower variance, but cost is $O(K)$ per *full model*, and you need K large before it helps on models with many latents.
- What if we could afford K samples *per latent variable*, not per whole vector?
Naively that's $K^{\#\text{latents}}$ combinations to consider: looks hopeless.

This is the problem Tensor Monte Carlo, then Massively-Parallel RWS, and then the source term trick, solve in turn.

Wake-Sleep, briefly

Classic wake-sleep (Hinton et al.) trains a generative model P and a recognition/proposal model Q together:



Reweighted Wake-Sleep (RWS) improves the Q update: instead of training Q on P 's own samples, use self-normalized importance weights from K samples of Q to push Q toward the *true posterior* directly. Unlike VI, RWS never reparameterizes: it works directly with **discrete latent variables**, no REINFORCE or relaxations needed.


Tensor Monte Carlo

The key move that makes “ K samples per latent” tractable:

- Draw K candidate values for *each* latent variable separately, not K whole-vector samples.
- Most latents only interact with a handful of neighbours in the model’s dependency graph (plates, conditional independence).
- So the $K^{\#\text{latents}}$ sum factorizes into a sequence of small tensor contractions (like einsum) that follow the graph structure: cost becomes polynomial in K , not exponential.


Why this matters: a 3-latent chain, $K = 3$

naive: all $K^3 = 27$ combinations



K slices of $K \times K$

structured: contract one edge at a time



$2K^2 = 18$ multiply-adds, not 27

$$\hat{P}_{\text{MP}} = \sum_{k_a, k_b, k_c} f_A(k_a) f_{AB}(k_a, k_b) f_{BC}(k_b, k_c) = \sum_{k_c} \left[\underbrace{\sum_{k_b} v_B[k_b] f_{BC}(k_b, k_c)}_{\text{contract } B} \right]$$

where $v_B[k_b] = \sum_{k_a} f_A(k_a) f_{AB}(k_a, k_b)$ is computed once and reused for every k_c .

What TMC gives you (and doesn't)

- + A much lower-variance, K -per-latent marginal likelihood estimator, at polynomial cost.
- On its own, it's just an estimator: no prescription for *training* P and Q together, and no cheap way to get posterior samples/moments back out afterward.

Massively-Parallel Reweighted Wake-Sleep (Heap et al., UAI 2023) turns this estimator into a full training algorithm.

Massively-Parallel RWS

The estimator: $P_{\text{MP}}(z)$

For each latent, draw K particles from Q ; contract following the plate structure (as above) to get a single scalar:

$$\hat{P}_{\text{MP}}(x) = (\text{TMC-style contraction of } p(z, x)/q(z) \text{ over all latents' } K \text{ particles})$$

- **Unbiased:** $\mathbb{E}[\hat{P}_{\text{MP}}(x)] = p(x)$, the true marginal likelihood, proved directly, not just assumed.
- This single fact is what later makes *exact* MCMC schemes (Particle MCMC) possible on top of this machinery; more on that at the end.

Two updates, one estimator

Practical form of the massively-parallel RWS updates (Heap et al., Eq. 19):


$$\Delta\theta_{\text{MP}} = \mathbb{E}_{Q_{\text{MP}}(z|x)} [\nabla_{\theta} \log \hat{P}_{\text{MP}}(z)] \quad \Delta\phi_{\text{MP}} = \mathbb{E}_{Q_{\text{MP}}(z|x)} [\nabla_{\phi} (-\log \hat{P}_{\text{MP}}(z))]$$

- **Wake phase:** ascend $\Delta\theta_{\text{MP}}$, i.e. climb $\log \hat{P}_{\text{MP}}$ w.r.t. P 's parameters.
- **Sleep phase:** ascend $\Delta\phi_{\text{MP}}$, i.e. *descend* $\log \hat{P}_{\text{MP}}$ w.r.t. Q 's parameters, which pushes Q toward the true posterior.

Equivalent to the fuller self-normalized-importance-weight form (Eq. 18): a weighted average of $\nabla \log P_{\theta}(z^{\mathbf{k}}, x) / \nabla \log Q_{\phi}(z^{\mathbf{k}}, x)$ over all K^n combinations, weighted by $r_{\mathbf{k}}(z) / \hat{P}_{\text{MP}}(z)$.

Both forms reuse the same $\hat{P}_{\text{MP}}(z)$ computation; no separate machinery for P vs. Q .

Results: MP RWS vs. global RWS, MovieLens



Predictive log-likelihood vs. $K \in \{3, 10, 30\}$, at three dataset sizes. MP RWS (red) beats global RWS (blue) at every K and every size, and the gap widens as the number of users grows (Heap et al., 2023).

MPIW

From “global” to massively-parallel importance weighting

MP-RWS trains P and Q . **MPIW** (Bowyer, Heap & Aitchison, UAI 2024) asks a different question: once trained, how do we get *posterior expectations* back out?

- **Global** importance weighting: draw K whole-vector samples, reweight, average: the classical scheme.
- **Massively-parallel** importance weighting (MPIW): draw K *per latent*, reweight *all* $K^{\#\text{latents}}$ combinations (same trick as before), and use *that* as the importance-sampling estimator of a posterior expectation.

This is the same $\hat{P}_{\text{MP}}(z)$ construction as before, now reframed as a general-purpose importance-weighting tool in its own right, not just something living inside the RWS training loop.

The math

With $\mathbf{k} = (k_1, \dots, k_n)$ indexing one combination of particles (one per latent), and $\text{qa}(i)$ the parents of latent i under Q :

$$r_{\mathbf{k}}(z) = \frac{P(x, z^{\mathbf{k}})}{\prod_i Q_{\text{MP}}(z_i^{k_i} \mid z_j \text{ for } j \in \text{qa}(i))} \quad \hat{P}_{\text{MP}}(z) = \frac{1}{K^n} \sum_{\mathbf{k} \in \mathcal{K}^n} r_{\mathbf{k}}(z)$$

and, for any moment function $m(\cdot)$, the posterior-expectation estimate:

$$\hat{m}_{\text{MP}}(z) = \frac{1}{K^n} \sum_{\mathbf{k} \in \mathcal{K}^n} \frac{r_{\mathbf{k}}(z)}{\hat{P}_{\text{MP}}(z)} m(z^{\mathbf{k}})$$

$r_{\mathbf{k}}(z)$ factorizes into low-rank tensors along the graph (as on the TMC slide), so both sums are computed via `opt_einsum`, never by looping over all K^n terms directly.

The shape of an MPIW call

```
z_i: [K] for each latent i      # one axis per latent
r_k = P(x, z^k) / Q_MP(z^k)    # tensor, shape [K]*n
P_MP = r_k.mean(all k axes)    # scalar
m_MP = ((r_k / P_MP) * m(z^k)).mean(all k axes)
```

- Input: K particles per latent. Output: a scalar ($\hat{P}_{\text{MP}}$) or a moment ($\hat{m}_{\text{MP}}$), never a K^n -sized object materialized in memory.
- Same shape whether you're computing the ELBO, an RWS update, or a posterior moment; only $m(\cdot)$ changes.

The source term trick

The remaining problem

MPIW gives an unbiased posterior-expectation *estimator*. In principle you plug in the quantity you want and reweight.

But once you have a trained P, Q : how do you get

- posterior **samples**,
- posterior **marginals** (the per-latent weight distribution over its K particles),
- posterior **moments** (means, variances, covariances)

back out, **without writing a bespoke resampling algorithm for every different plate topology?**

The trick

Using Autodiff to Estimate Posterior Moments, Marginals and Samples (Bowyer, Heap & Aitchison, UAI 2024):

1. Add an auxiliary **source term** J into the log joint, additively, before running the *exact same* K -dim contraction used to compute the bound.
2. That contraction now computes a scalar $L(J)$ instead of just $L(0)$.
3. Differentiate $L(J)$ w.r.t. J , at $J = 0$, using ordinary autodiff. The derivative is the marginal / moment you wanted.

Exactly the same move as getting moments from $\partial \log Z[J] / \partial J$ for a log-partition function with an external source J : a very old idea in statistical physics and exponential families, applied here to a Monte Carlo estimator instead of an analytic one.

Why this is the elegant part

- No new algorithm per quantity: `alan` computes posterior samples, marginals, *and* moments through literally the same code path that computes the bound.
- Works uniformly across arbitrarily nested plates: the contraction already handles that; the source term rides along for free.
- It's autodiff doing the bookkeeping that would otherwise be a hand-written, model-specific resampling routine.

This is really in the code

From alan's `Sample.py` (lightly trimmed):

```
J_tensor = zeros(shape, requires_grad=True)
L = self._elbo(sample, extra_log_factors={...: J_tensor})
marginals = grad(L, J_tensor_list)
```

- `J_tensor` starts at all zeros: it doesn't change what `L` evaluates to.
- But differentiating `L` through it gives back exact per-particle marginal weights, for whichever latent(s) `J` was attached to.

QEM

Training Q still needs gradients

MP-VI and MP-RWS both update Q 's parameters ϕ by **gradient ascent** on a massively-parallel bound.

- Needs a learning rate, tuned per model.
- **Not invariant to reparameterization**: rescale a latent by a constant (leaving the actual distribution over data untouched) and gradient descent's behaviour changes, and sometimes it stops converging at all.

QEM (Heap, Bowyer & Aitchison, 2025) removes gradient descent from the Q update entirely.

The QEM idea

E-step

Use MPIW + the source term trick to compute the *current* posterior moments, exactly as before.

M-step


Set Q 's *next* parameters directly from those moments, in closed form, for exponential-family Q (Gaussian, Gamma, Beta, Dirichlet, Binomial, ...).

The name: it resembles EM, but optimizes the parameters of the *approximate posterior* Q , not the model's prior/likelihood parameters θ , as classical EM does. Hence "QEM."

```
for t in 1..T:  
    Q_t = set_exp_family_params(m[t-1]) # M-step  
    z = sample(Q_t)  
    m_new = MPIW_moments(z, Q_t)      # E-step  
    m[t] = lambda*m_new + (1-lambda)*m[t-1] # smoothing
```

- No gradients, no learning rate: just moment matching.
- The exponential moving average on the last line matters: a single noisy moment estimate can otherwise cause pathological collapse (e.g. near-zero effective weight on all but one sample). Provably converges to the unbiased, zero-variance estimate as $T \rightarrow \infty$.

Results: five real hierarchical models, ELBO vs. time




QEM (pink) reaches the same ELBO / predictive log-likelihood as MP-RWS (green) and MP-VI (orange) in a fraction of the wall-clock time, on Bus Breakdown, MovieLens100K, Occupancy, Radon, and Covid (Heap, Bowyer & Aitchison, 2025). No backprop through the whole joint per step.

Also: exactly reparameterization invariant

- **At least as accurate:** posterior first-moment MSE (vs. long HMC runs) at least as good as MP-RWS, and clearly better than MP-VI, throughout.
- **Reparameterization invariant, exactly:** rescale a latent by $\times 1/100$ (then back), and QEM's trajectory is unchanged. MP-VI *fails to converge at all within 250 iterations* on the reparameterized Radon and Covid models: a real, measured failure, not a hypothetical one.

How it fits together

The whole picture



Further work

Further work: ideas building on $\hat{P}_{\text{MP}}$'s unbiasedness

- **How biased is MP-IS in practice, and does it actually close the gap to HMC?** An empirical open question worth measuring directly on a model with a known answer.
- **Better debiasing than jackknife:** chain-recycling schemes (BR-SNIS-style) that target the already plate-reduced estimator instead of the raw joint.
- **Exact inference via Particle Gibbs:** $\hat{P}_{\text{MP}}$'s unbiasedness is exactly the ingredient Particle MCMC needs; pin a reference trajectory and get an exact posterior, not an approximate one.
- **Ancestor sampling (PGAS)** as the likely key to making Particle Gibbs mix well, given the plate-wise ESS unevenness the source term trick already lets us measure.

Wrap-up

Summary

- TMC makes K -per-latent tractable by contracting along the plate graph.
- MP-RWS turns that estimator into a training algorithm, and proves $\hat{P}_{MP}$ unbiased.
- MPIW reframes the same estimator as a general tool for posterior expectations, not just RWS training.
- The source term trick gets samples/marginals/moments out of the *same* computation, via autodiff, for free.
- QEM closes the loop: uses those moments as a gradient-free, reparameterization-invariant E-step, beating MP-VI and MP-RWS on both speed and (at worst) matching them on accuracy.
- That whole stack opens up further work: measuring MP-IS bias in practice, better debiasing, and exact inference via Particle Gibbs, all building directly on $\hat{P}_{MP}$'s unbiasedness.

Questions